


12-T3: GCPS REAL-TIME BUS MONITORING SYSTEM

CAPSTONE INDUSTRY PARTNER PROJECT

CS 4850 - SECTION 01 – FALL 2024

PROFESSOR PERRY – NOVEMBER 13, 2024




Sam Bostian
(Team Leader)




Michael Rizig
(Developer)



Charlie McLarty
(Developer)



Brian Pruitt
(Documentation)



Allen Roman
(Developer)

Name	Role	Cell Phone / Email
Sam Bostian	Team Leader	321.292.4693 Bostian.sam@gmail.com
Michael Rizig	Developer	678.668.3294 mrizig@students.kennesaw.edu
Charlie McLarty	Developer/QA	470.303.9544 cmclarty21@gmail.com
Brian Pruitt	Documentation	404.207.6548 bpruitt9@students.kennesaw.edu
Allen Roman	Developer	470.249.0421 aroman14@students.kennesaw.edu

Website:	http://12-t3-gwinnett-bus.infinityfreeapp.com/?i=1
Github:	https://github.com/Team3-GwinnettBus/Gwinnett-Bus

Lines of Code	Over 2.5M!
Number of Components/Tools	6 (<i>Kafka, SQL Server, Podman, React.js, Flask, Leaflet</i>)
Total Man Hours	Approx. 1250 Hours

TABLE OF CONTENTS

Introduction	4
Project Overview	4
Project Objectives	4
Scope	4
Requirements	5
Functional requirements	5
Non-Functional Requirements	5
Constraints	5
Design	6
System Architecture	6
Data Flow	6
Containerization	6
Development	6
Development Approach	6
System Components	6
Containerization	7
Test Report	8
Version Control	9
Summary	10
Project Achievements	10
Lessons Learned	10
Future recommendations	11
Appendix	11
Project Plan	11
Gantt Chart	11
Project Set-Up	12
Install Required Tools	12
Configure Environment	12
Build and Deploy Containers	12

Monitor and Test	13
Shut Down and Clean Up	13
Server Login	14
How to login to the project server:	14
Useful Linux Commands for Groups and Users	16
How to Create a Group.....	16
How to Create a User	16
How to Add a Password to a User	17
How to Add a User a Group.....	17
Software Package Installations	17
How to Upgrade Packages.....	17
How to Install Python 3.....	18
Python Packages to Install.....	19
How to Install Java	20
How to Install Podman	20
How to Install Zookeeper.....	21

INTRODUCTION

PROJECT OVERVIEW

The GCPS Real-Time Bus Monitoring System was developed to enhance the efficiency of monitoring bus operations for Gwinnett County Public Schools. The project replaces an outdated polling-based system with Kafka-based real-time event streaming, enabling GCPS to access real-time telemetry data for over 2,000 buses.

PROJECT OBJECTIVES

The key objectives of this project are:

1. **Reduce Latency:** Switch from a polling system to Kafka event streaming to improve data freshness and reduce retrieval delays.
2. **Enable Scalable Deployment:** Utilize Podman containerization for easy deployment and scalability across environments.
3. **Improve Data Processing and Storage:** Validate and store bus telemetry data in an SQL Server database for consistent, reliable access.

SCOPE

The scope of this project includes data ingestion, validation, storage, and visualization of bus telemetry data, as well as establishing a deployment pipeline using Podman for container management.

REQUIREMENTS

FUNCTIONAL REQUIREMENTS

Real-Time Data Ingestion: The system must consume real-time telemetry data from Kafka, capturing information like bus location, speed, and heading.

Data Validation: Incoming data must be validated for accuracy before storage.

Data Storage in SQL Server: Validated data should be stored in SQL Server for easy access.

Frontend Visualization: The system should provide a visual map showing bus locations, updated in near real-time.

Deployment Automation: The system must be deployable through Podman containers, ensuring consistent setup across environments.

NON-FUNCTIONAL REQUIREMENTS

Performance: The system must process and display data with minimal latency.

Scalability: The system must handle telemetry data from up to 2,000 buses.

Security: Data should be encrypted and secured, with access limited to authorized users.

CONSTRAINTS

- The system must run on Red Hat Linux environments.
- GCPS is a Microsoft shop and uses Microsoft SQL Server.
- The entire solution should be containerized through Podman.

DESIGN

SYSTEM ARCHITECTURE

The system consists of four main components:

1. **Kafka Producer:** Simulates and streams bus telemetry data to Kafka.
2. **Kafka Consumer:** Consumes, validates, and stores data in SQL Server.
3. **SQL Server:** Central storage for validated telemetry data.
4. **Frontend:** A React.js application that visualizes bus locations on a map.

DATA FLOW

Producer generates simulated telemetry data and sends it to Kafka.

Consumer receives the data, performs validation, and stores it in SQL Server.

Frontend queries SQL Server periodically to update the map with the latest bus locations.

CONTAINERIZATION

Using Podman, each component (producer, consumer, frontend) is containerized for consistent and scalable deployment. The pod.yml file configures the network and connects containers within the same pod, enabling seamless communication.

DEVELOPMENT

DEVELOPMENT APPROACH

We used an Agile approach with sprints focused on building core components like the Kafka producer, Kafka consumer, and frontend. Iterative feedback allowed us to refine requirements and implementation details to meet sponsor needs.

SYSTEM COMPONENTS

Kafka Producer: Simulates real-time data and streams it to Kafka.

Kafka Consumer: Validates and stores telemetry data in SQL Server.

Frontend Visualization: Uses React.js and Leaflet to display bus locations.

CONTAINERIZATION

Each component was containerized using Podman, enabling easy deployment and scalability. Docker files were created for each component, specifying dependencies and startup commands.

TEST REPORT

Requirement	Description	Pass	Fail	Severity
Test Requirement	This is a Test Requirement to serve as an example	✓	✗	Low
Real-time data ingestion through Kafka	Verify that the Kafka producer sends data in real time	✓		High
Data validation rules	Check that invalid GPS data is flagged and logged	✓		High
Data storage in SQL Server	Confirm that validated data is stored accurately	✓		High
Error handling for invalid data	Verify that invalid data is stored in separate table	✓		Medium
Container initialization	Check if all containers start without errors	✓		High
Inter-container communication	Confirm Kafka-to-SQL data flow in pod environment	✓		High
Pod shutdown and cleanup	Ensure stop_pod.sh script removes all containers/pods	✓		Medium
Monitoring setup (optional feature)	Verify monitoring tool displays container status	✓		Low
Scalability test with 2,000 bus data points	Check system performance under high data volume	✓		High
System latency under load	Measure if latency remains within acceptable range	✓		High
Data retrieval from SQL Server	Verify that bus location data is retrievable	✓		High
Consistent container deployment across environments	Test pod deployment on different environments	✓		Medium

VERSION CONTROL

Version	Primary Author	Description of Version	Changes	Date Completed
0.1	Sam Bostian, Charlie McLarty, Brian Pruitt, Michael Rizig	Initial Release	Initial Project Plan	
1.0	Michael Rizig Sam Bostian	Initial server and Kafka setup.	<ul style="list-style-type: none"> • RHEL 9.4 server software install • Kafka and Database Setup 	Sept. 1, 2024
1.1	Michael Rizig	Fix the connection error with Microsoft SQL.	<ul style="list-style-type: none"> • Reinstalled MSSQL with unixODBC driver 	Sept. 23, 2024
2.0	Charlie McLarty	Initial telemetry data simulation .	<ul style="list-style-type: none"> • A Kafka producer script was written to simulate GCPS Buses telemetry data 	Sept. 28, 2024
2.1	Charlie McLarty	Simulated buses were added to the front end and move along a path.	<ul style="list-style-type: none"> • Random buses were generated to fit inside the boundaries of Gwinnett County 	Oct. 4, 2024
3.0	Charlie McLarty	Turned Gwinnett County into a connected graph.	<ul style="list-style-type: none"> • Gwinnett County was transformed into a connected graph so Buses can travel along paths that overlap the roads in Gwinnett County 	Oct. 5, 2024

3.1	Charlie McLarty	Geofences added	<ul style="list-style-type: none"> A list of coordinates for all GCPS schools was created and geofences were created around the schools. 	Oct. 9, 2024
4.0	Allen Roman	Project Containerized	<ul style="list-style-type: none"> A docker file was created and the project was containerized for easy system install. 	Nov. 1, 2024
4.1	Allen Roman	Shell Commands added	<ul style="list-style-type: none"> Shell commands were added to allow for easy deployment of all containers. 	Nov. 3, 2024
4.2	Allen Roman	Switched base image	<ul style="list-style-type: none"> The base image for the docker file was switched from python:3.10-slim to rhel-ubi8-python 	Nove. 12, 2024

SUMMARY

PROJECT ACHIEVEMENTS

- Real-time data ingestion and validation through Kafka.
- Scalable, containerized deployment with Podman.
- Interactive frontend displaying real-time bus locations.

LESSONS LEARNED

Agile development allowed us to adapt quickly to changing requirements.

Containerization with Podman provided a robust and consistent deployment environment.

FUTURE RECOMMENDATIONS

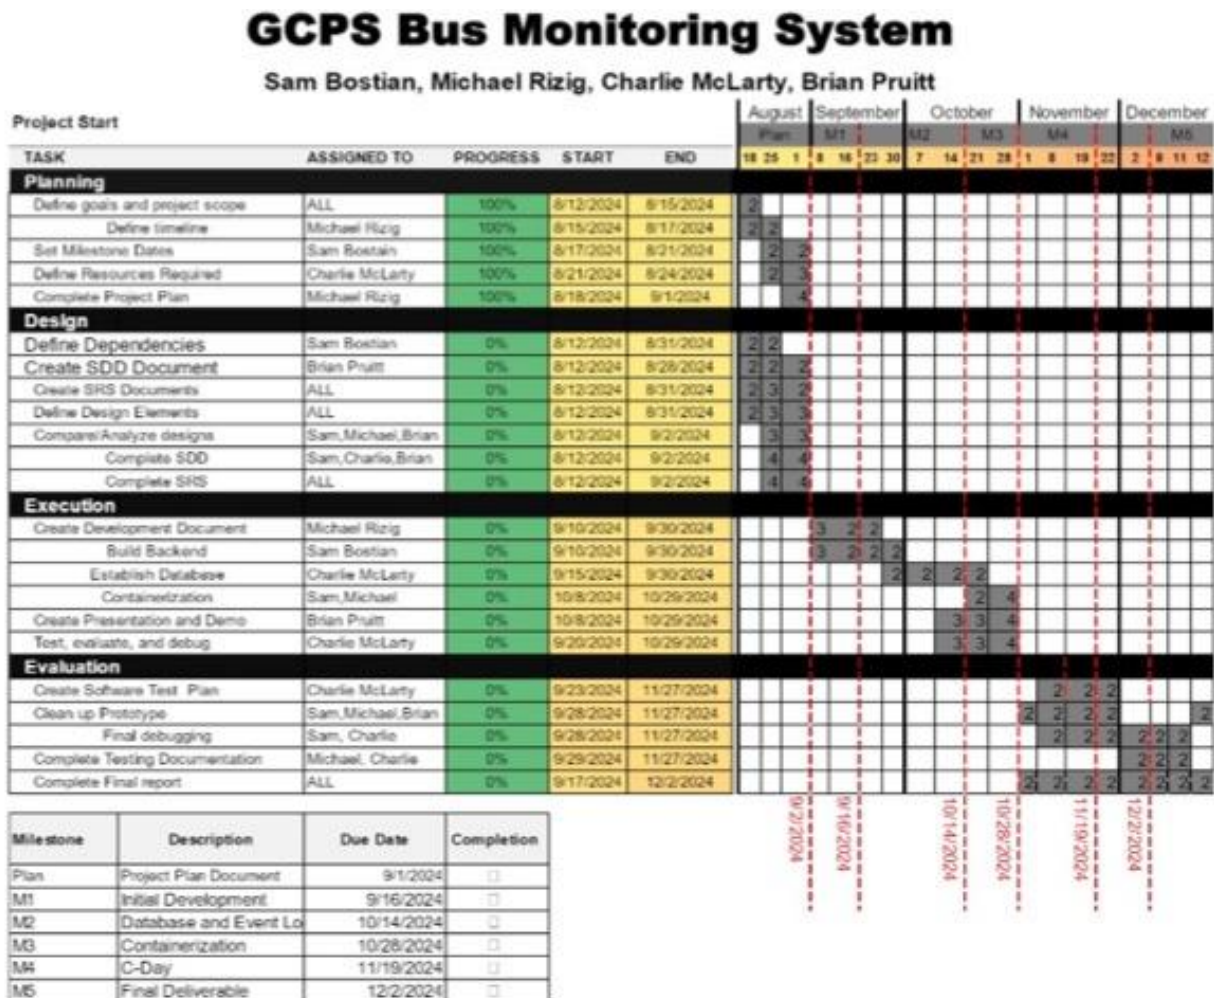
Implement advanced monitoring for better system insights.

Extend the system to support additional data sources or integrate with other school district systems.

APPENDIX

PROJECT PLAN

GANTT CHART



PROJECT SET-UP

INSTALL REQUIRED TOOLS

Podman: Install Podman for containerization. This tool will be used to create, manage, and deploy containers for the project's components.

- Command: On a Red Hat-based system, use **sudo dnf install podman**.

Kafka: Install and configure Kafka for event streaming.

- Ensure that Kafka is correctly configured to create and manage topics that will handle bus telemetry data.

SQL Server: Install and configure Microsoft SQL Server, as it will store validated telemetry data from Kafka.

- Configure SQL Server to allow connections from the Podman containers and set up necessary user credentials.

CONFIGURE ENVIRONMENT

Environment Variables: Set up environment variables as defined in the README.md file.

These include:

- **SERVER_IP:** IP address of the server hosting Kafka and SQL Server.
- **KAFKA_TOPIC:** The name of the Kafka topic to which the producer will send data.
- **DB_NAME, DB_USER, DB_PASSWORD:** Database name and credentials for SQL Server.

Verify that these environment variables are accessible to each container by configuring them in the pod's YAML file or exporting them before deployment.

BUILD AND DEPLOY CONTAINERS

Use the **run_pod.sh** script to build and deploy the containers within a Podman pod.

This script will:

- Build the **producer container** that simulates bus data and sends it to Kafka.
- Build the **consumer container** that consumes data from Kafka and interacts with SQL Server.

- Start these containers within a shared Podman pod network, ensuring they can communicate.

Run **run_pod.sh** from the terminal to initiate this setup

MONITOR AND TEST

Container Status: Confirm that each container is running as expected by checking Podman's status output.

- Command: **podman ps** to view running containers.
- Command: **podman logs <id>** to view logs for running container.

Data Flow Verification:

- Verify that the Kafka producer is successfully sending data to the Kafka topic.
- Confirm that the Kafka consumer is receiving this data and inserting it into SQL Server.
- Run sample API requests to ensure the data can be retrieved and displayed accurately in the web interface.

SHUT DOWN AND CLEAN UP

When finished, use the stop_pod.sh script to stop and remove all containers and pods, freeing up resources and leaving a clean environment for future deployments.

Run **stop_pod.sh** from the terminal:


- This script will:
 - a. Stop and remove all containers and pods associated with the project.
 - b. Prune any unused resources, such as images, to prevent storage clutter.

Server Login

HOW TO LOGIN TO THE PROJECT SERVER:

1. If Kennesaw State's VPN is not installed follow these instructions:
 - a. Instructions can be found here: <https://uits.kennesaw.edu/vpn/index.php>
2. If windows Linux subsystem(WSL) is not installed follow these instructions:
 - a. Instructions can be found here:
<https://techcommunity.microsoft.com/t5/windows-11/how-to-install-the-linux-windows-subsystem-in-windows-11/m-p/2701207/page/3>
 - b. It is suggested to upgrade to WSL2 using the following instructions if your system is currently using WSL: <https://learn.microsoft.com/en-us/windows/wsl/install>
 - i. **Scroll down to the section titled Upgraded version from WSL 1 to WSL 2:**
<https://learn.microsoft.com/en-us/windows/wsl/install#upgrade-version-from-wsl-1-to-wsl-2>
 - ii. It is recommended to pin the WSL to the start or task bar.

3. You must be logged into the KSU VPN and in the Portal: vpn.kennesaw.edu to successful execute the following steps (the world must be colored and not in grayscale):



- a. Open the WSL and type the following command lines to access the server provided by the school:
`ssh [username]@[ip address]`
- b. The user will be prompted to enter a password.
- c. Congratulations, you are logged in.

USEFUL LINUX COMMANDS FOR GROUPS AND USERS

In Linux a group is an assembly of users. An administrator might want to create a group in Linux to simplify what users have access to specific files. Groups can be created from the root folder or users that have superuser do (**sudo**) access.

HOW TO CREATE A GROUP

```
$ sudo groupadd [group name]
```

-- Or for more control to set the group ID(GID)--

```
$ sudo groupadd -g [GID 1000-60000] [group name]
```

HOW TO CREATE A USER

```
$ sudo useradd {options} [username]
```

--Example of an {option}: adding a custom user ID(UID)--

```
$ sudo useradd -u [UID 1000-60000] [username]
```

--Or to add a String to identify the user like their name--

```
$ sudo useradd -c "[String]" [username]
```

HOW TO ADD A PASSWORD TO A USER

```
$ sudo passwd [username]
```

HOW TO ADD A USER A GROUP

```
$ sudo usermod --append --groups [group name] [username]
```

Software Package Installations

HOW TO UPGRADE PACKAGES

--Upgrade All packages--

```
$ sudo dnf upgrade
```

--Upgrade a Single Package--

```
$ sudo dnf upgrade [package name]
```

--Upgrade a Specific Package Group--

```
$ sudo dnf upgrade [group name]
```

Select one of the following Python versions:

--Install Python 3.9--

```
$ sudo dnf install python3
```

--Install Python 3.11--

```
$ sudo dnf install python3.11
```

--Install Python 3.12--

```
$ sudo dnf install python3.12
```

--Install Python Request Modules--

```
$ sudo dnf install python3-requests
```

--Install Python3 package manager(pip)--

```
$ sudo dnf install python3-pip
```

--Install Python3.11 package manager(pip)--

```
$ sudo dnf install python3.11-pip
```

--Install Python3.12 package manager(pip)--

```
$ sudo dnf install python3.12-pip
```

--Install Additional Python3Developer tools--

Enable the following subscription manager code before the following code

```
$ sudo subscription-manager repos --enable codeready-builder-for-rhel-9-x86_64-rpms
```

Then install the following development tools for python:

- ☐ \$ **sudo** dnf install python3*-idle
- ☐ \$ **sudo** dnf install python3*-debug
- ☐ \$ **sudo** dnf install python3*-Cython
- ☐ \$ **sudo** dnf install python3.11-pytest
- ☐ \$ **sudo** dnf install python3.12-pytest
- ☐ \$ **sudo** dnf install python3.12-pytest
- ☐ \$ **sudo** pip3 install Flask

HOW TO INSTALL JAVA

Select one of the following OpenJDKs versions:

--Install OpenJDK11--

```
$ sudo dnf install java-11-openjdk
```

--Install OpenJDK17--

```
$ sudo dnf install java-17-openjdk
```

--Install OpenJDK8--

```
$ sudo dnf install java-1.8.0-openjdk
```

HOW TO INSTALL PODMAN

```
$ sudo dnf install podman
```

****Warning ZooKeeper was marked as deprecated in Apache Kafka with the release of version 3.5 and will be completely removed with the Kafka 4.0 release****

HOW TO INSTALL ZOOKEEPER

Complete the following steps in order to install zookeeper

(1) Install the Extra Packages for Enterprise Linux (EPEL) repository onto the system

```
$ sudo dnf install https://dl.fedoraproject.org/pub/epel/epel-release-latest-9.noarch.rpm  
$ sudo dnf upgrade
```

(2) Add the following extra repositories

```
$ sudo subscription-manager repos --enable "rhel-* -optional-rpms" --enable "rhel-* -  
extras-rpms"  
$ sudo dnf update
```

(3) Install Snap

```
$ sudo dnf install snapd
```

(4) Enable the systemd unit that manages the main snap communication socket

```
$ sudo dnf systemctl enable --now snapd.socket
```

(5) Enable the classic Snap support

```
$ sudo ln -s /var/lib/snapd/snap /snap
```

(6) Install Zookeeper

```
$ sudo snap install zookeeper
```

